

Executive Summary

Overview

We are building the first **secure, blockchain-based AI oversight platform** to ensure **trustworthy, verifiable AI agent interactions**. Our system prevents **collusion, hallucinations, and misalignment**, enabling AI agents to operate transparently and in alignment with stakeholders.

Problem & Opportunity

AI-driven automation lacks **real-time oversight**, leading to **hallucinations, misalignment, and opaque decision-making**. Businesses and regulators demand **greater trust, accountability, and intervention mechanisms** for AI operations.

Solution & Differentiation

Unlike static AI compliance tools, we provide:

- **Live AI interaction monitoring** to detect misalignments in real time.
- **Multi-layer verification** using LLMs, formal logic, and expert models.
- **Human-aligned interpretability** for transparent AI decisions.
- **Intervention mechanisms** to prevent AI failures.
- **Blockchain-backed security** ensuring tamper-proof oversight.

Market & Business Model

We target **B2C enterprises**, initially focusing on **consumer goods companies and small businesses** looking to **reduce customer service costs**. Our **subscription-based SaaS model** scales with verification depth and enterprise needs.

Team & Vision

Ming-Chang Chiu (Founder & CEO)

- **PhD in Computer Science (2025), USC | Ex-Google Research, ex-NVIDIA**
- **ICML 2024 Best Paper Award | Expert in Trustworthy AI & Agents**
- Soon to recruit a **technical co-founder** for scaling.

Why Now?

AI adoption is accelerating, but **trust and oversight remain unsolved challenges**. Our platform sets a **new standard for AI governance**, making autonomous AI safer, smarter, and aligned with human intent.